

XIANGYU ZHOU

+01-2019120046 xiangyu@wayne.edu Google Scholar Xiangyu Zhou xzhou98 Website

RESEARCH INTERESTS

Trustworthy AI & Large Language Model & Large Reasoning Model & Generative Retrieval & Agentic AI

EDUCATION

Wayne State University, Detroit, Michigan, USA.

Aug 2023 – Present

Ph.D. in Computer Science

Advisor: Dr. Dongxiao Zhu

Stevens Institute of Technology, Hoboken, New Jersey, USA.

Aug 2021 – May 2023

M.S. in Computer Science

Chongqing University of Posts and Telecommunications, Chongqing, China.

Sep 2017 – May 2021

B.S. in Software Engineering

ONGOING RESEARCH PROJECTS

- **Safety and Helpfulness Alignment of LLMs:** Building reinforcement-learning and contrastive-learning pipelines to improve model safety and helpfulness, including custom safety/helpfulness datasets, format rewards for GRPO, and post-training fine-tuning of LLMs.
- **Robustness of LRMs under In-Context Learning:** Studying how frontier reasoning models fail under multi-turns or few-shots conversations and developing post-training methods to improve robustness without degrading general performance.

WORK EXPERIENCE

Trustworthy AI Lab, Wayne State University

Aug 2023 – Present

Graduate Research Assistant

Detroit, MI

- Conduct research on trustworthy LLMs, focusing on unlearning, safety alignment, and reasoning robustness.

LinkedIn Corporation

May 2026 – Aug 2026

AI/ML Engineer - Generative AI, Intern

Mountain View, CA

- Developed **generative AI models for candidate retrieval** in LinkedIn's hiring assistant, improving hiring recall by **46.4%** and role fit by **33.3%** in offline evaluation through reinforcement learning.
- Built the team's **first reinforcement learning pipeline** for candidate retrieval using LLM-based rewards. Created a **candidate-profile lookup table** to replace online requests and reduce training latency.
- Implemented **constrained decoding** to prioritize candidates actively seeking jobs, increasing their share among retrieved candidates.

Shandong Googosoftware Co., Ltd.

Jul 2019 – Aug 2019

Full Stack Developer, Intern

Shandong, China

- Contributed to the development of a university Asset Management Platform used for financial tracking and asset reporting.
- Improved front-end reliability of web app by debugging JavaScript interaction issues (e.g., unresponsive UI elements).

PUBLICATIONS

Peer-Reviewed Publications

- Zhou, X., et al. "Not All Tokens Are Meant to Be Forgotten." *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 40. No. 44. 2026. (Accepted as Oral, AAAI-26).
- Zhou, X., et al. "Hijacking large language models via adversarial in-context learning." *European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in Databases (ECML-PKDD -26)*.

- Zade, S. Z., **Zhou, X.**, Liu, S., Zhu, D. "Attention Smoothing Is All You Need For Unlearning" *The Fourteenth International Conference on Learning Representations (ICLR-26)*.
- Li, M., **Zhou, X.**, Wu, J. et al. "Unraveling salt-responsive genes in *Suaeda salsa* through genomic and transcriptomic profiling across salinity gradients" *BMC Genomics* 27, 475 (2026).
- Roshani, M. A., **Zhou, X.**, Qiang, Y., Suresh, S., Hicks, S., Sethuraman, U., & Zhu, D. "Generative large language model—powered conversational ai app for personalized risk assessment: Case study in covid-19." *Published in the Journal of Medical Internet Research (JMIR)*.
- Sultan, R., Zhu, H., **Zhou, X.**, Li, C., Khanduri, P., Brocanelli, M., Zhu, D., et al. "WalkGPT: Grounded Vision–Language Conversation with Depth-Aware Segmentation for Pedestrian Navigation." *The IEEE/CVF Conference on Computer Vision and Pattern Recognition 2026 (CVPR-26)*.
- Zade, S. Z., Qiang, Y., **Zhou, X.**, Zhu, H., Roshani, M. A., Khanduri, P., & Zhu, D. "Automatic Calibration for Membership Inference Attack on Large Language Models" *In Proceedings of the 28th European Conference on Artificial Intelligence (ECAI-25)*.
- Zheng, W., Walquist, E., Datey, I., **Zhou, X.**, Berishaj, K., ... & Zytco, D. "It's Not What We Were Trying to Get At, but I Think Maybe It Should Be": Learning How to Do Trauma-Informed Design with a Data Donation Platform for Online Dating Sexual Violence. *In Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems (CHI-24)*.
- Zheng, W., Walquist, E., Datey, I., **Zhou, X.**, Berishaj, K., McDonald, M., ... & Zytco, D. "Towards trauma-informed data donation of sexual experience in online dating to improve sexual risk detection AI." *In proceedings of the 36th annual ACM symposium on user interface software and technology (UIST-23)*.
- Walquist, E., Datey, I., Zheng, W., **Zhou, X.**, Berishaj, K., McDonald, M., ... & Zytco, D. "Collective Consent: Who Needs to Consent to the Donation of Data Representing Multiple People?". *Proceedings of the ACM on Human-Computer Interaction (CSCW-25)*.

Preprints and Under Review

- **Zhou, X.**, et al. "Alignment of LLMs via Counter-Aligned Few-Shot Conversation Exposure." *Under review at The Fortieth Annual Conference on Neural Information Processing Systems (NeurIPS-26)*.
- **Zhou, X.**, et al. "Learning to Poison Large Language Models for Downstream Manipulation." *Under review at the European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in Databases (ECML-PKDD -26) arXiv:2402.13459*.

REVIEWING EXPERIENCE

Reviewer: AAAI Conference on Artificial Intelligence (**AAAI 2026, 2027**), International Conference on Machine Learning (**ICML 2026**), Conference on Neural Information Processing Systems (**NeurIPS 2026**), and International Conference on Learning Representations (**ICLR 2027**).

PROJECTS EXPERIENCE

Codebook Quality Evaluation Pipeline for GR Model (LinkedIn)

May 2026 – Present

AI/ML Engineer

- Built an end-to-end evaluation pipeline assessing training codebook quality for Generative Recruiting models; instrumented data checks, drift detection, and robustness analyses across cohorts and time.
- Proposed a **new task-specific metric** aligned to LinkedIn hiring downstream objectives (match relevance, conversion propensity), enabling actionable benchmarking and guiding codebook improvements.

Ube (Data Donation Android Application)

Feb 2023 – Apr 2024

Full Stack Developer

- Independently developed an Android application with **React-Native** as the framework, and **Firestore** as the database.
- **Collaborated in writing and co-authoring** the research paper detailing the application's architecture and its implications for data privacy.
- The papers were successfully published at UIST-23, CHI-24, and CSCW-25, contributing to the discourse on ethical data usage and privacy technologies.

Movies Recommendation System

Oct 2021 – Dec 2021

Full Stack Developer

- Developed a movie communication website with **React** as the front end, **Node.js** as the back end, and **MongoDB** as the database.
- Implemented login, sign up and log out, finished hashing password and matching to password with **Bcrypt**, and improved **security** for user login verification.
- Accomplished insertion, updating, deletion, and querying of information of user and movie's comment, imported **XSS** defense to protect the system from Cross-Site Scripting Attack.

Campus Delivery System

Nov 2018 – Mar 2019

Front-end Developer

- Developed a web application with **HTML**, **JavaScript**, **jQuery**, and **Vue** as the front end.
- Designed and coded functions for items adding, querying, modifying, and deleting of goods and riders.

SKILLS

- **Programming:** Python, Java, Javascript
- **Frameworks:** PyTorch, HuggingFace, React, React-native, Vue
- **Tool:** Git, Firebase, MySQL, MongoDB
- **Machine Learning:** Agentic AI, Large Language Models, Large Reasoning Models, Trustworthy AI, Fine tuning, Reinforcement Learning